

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

#from google.colab import files
#uploaded = files.upload()

import pandas as pd
import io
#data = pd.read_csv(io.BytesIO(uploaded['housing.csv']))
data = pd.read_csv('/content/sample_data/housing.csv')
data.info()
```

```
[>] <class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms        20433 non-null  float64
5   population             20640 non-null  float64
6   households             20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
data.head()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population
0	-122.23	37.88	41.0	880.0	129.0	322.0
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0
2	-122.24	37.85	52.0	1467.0	190.0	496.0
3	-122.25	37.85	52.0	1274.0	235.0	558.0
4	-122.25	37.85	52.0	1627.0	280.0	565.0

Encoding

(1)Label encoder

(2)Onehot encoder

```
from sklearn.preprocessing import LabelEncoder , OneHotEncoder
data['median_house_value'].value_counts()
```

```

500001.0    965
137500.0    122
162500.0    117
112500.0    103
187500.0     93
...
303200.0     1
307900.0     1
383200.0     1
360800.0     1
405500.0     1
Name: median_house_value, Length: 3842, dtype: int64

```

```

le=LabelEncoder()
data['median_house_value']=le.fit_transform(data['median_house_value'])
data['median_house_value'].value_counts()

```

```

3841    965
959    122
1209    117
710    103
1459    93
...
3172     1
3275     1
3204     1
3091     1
2119     1
Name: median_house_value, Length: 3842, dtype: int64

```

```
le.classes_
```

```
array([ 14999.,  17500.,  22500., ..., 499100., 500000., 500001.])
```

(2)Onehot Encoder

```
data['ocean_proximity'].value_counts()
```

```

<1H OCEAN    9136
INLAND       6551
NEAR OCEAN   2658
NEAR BAY     2290
ISLAND        5
Name: ocean_proximity, dtype: int64

```

```

one_hot = OneHotEncoder()
transformed_data = one_hot.fit_transform(data['ocean_proximity'].values.reshape(-1,1)).toarray()
one_hot.categories_

```

```
[array(['<1H OCEAN', 'INLAND', 'ISLAND', 'NEAR BAY', 'NEAR OCEAN'],
      dtype=object)]
```

```

transformed_data = pd.DataFrame(transformed_data ,
                                columns = ['<1H OCEAN', 'INLAND', 'ISLAND', 'NEAR BAY', 'NEAR OCEAN'])

```

```
transformed_data.head()
```

	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN
0	0.0	0.0	0.0	1.0	0.0
1	0.0	0.0	0.0	1.0	0.0
2	0.0	0.0	0.0	1.0	0.0
3	0.0	0.0	0.0	1.0	0.0
4	0.0	0.0	0.0	1.0	0.0

```
transformed_data.iloc[90, ]
```

```
<1H OCEAN    0.0
INLAND       0.0
ISLAND       0.0
NEAR BAY     1.0
NEAR OCEAN   0.0
Name: 90, dtype: float64
```

```
data['median_house_value'][90]
```

```
1209
```

Normalization & Standardization

```
# consider only numerical columns
```

```
numeric_columns = [c for c in data.columns if data[c].dtype != np.dtype('O')]
numeric_columns
```

```
['longitude',
 'latitude',
 'housing_median_age',
 'total_rooms',
 'total_bedrooms',
 'population',
 'households',
 'median_income',
 'median_house_value']
```

```
len(numeric_columns) , len(data.columns)
```

```
(9, 10)
```

```
numeric_columns.remove('longitude')
numeric_columns.remove('latitude')
```

```
temp_data = data[numeric_columns]
temp_data
```

	housing_median_age	total_rooms	total_bedrooms	population	households	medi
0	41.0	880.0	129.0	322.0	126.0	
1	21.0	7099.0	1106.0	2401.0	1138.0	
2	52.0	1467.0	190.0	496.0	177.0	
3	52.0	1274.0	235.0	558.0	219.0	
4	52.0	1627.0	280.0	565.0	259.0	
...
20635	25.0	1665.0	374.0	845.0	330.0	
20636	18.0	697.0	150.0	356.0	114.0	
20637	17.0	2254.0	485.0	1007.0	433.0	
20638	18.0	1860.0	409.0	741.0	349.0	
20639	16.0	2785.0	616.0	1387.0	530.0	

20640 rows × 7 columns

Normalization

```
from sklearn.preprocessing import StandardScaler , MinMaxScaler
import warnings
warnings.filterwarnings('ignore')
normalizer = MinMaxScaler()
temp_data.dropna(axis = 1 , inplace = True)
normalized_data = normalizer.fit_transform(temp_data)
pd.DataFrame(normalized_data , columns = temp_data.columns)
```

	housing_median_age	total_rooms	population	households	median_income	median_house_value
--	--------------------	-------------	------------	------------	---------------	--------------------

Standardization

1	0.392157	0.180503	0.067210	0.186976	0.538027
---	----------	----------	----------	----------	----------

```
standard_scaler = StandardScaler()
standardized_data = standard_scaler.fit_transform(temp_data)
pd.DataFrame(standardized_data , columns = temp_data.columns)
```

	housing_median_age	total_rooms	population	households	median_income	median_house_value
0	0.982143	-0.804819	-0.974429	-0.977033	2.344766	
1	-0.607019	2.045890	0.861439	1.669961	2.332238	
2	1.856182	-0.535746	-0.820777	-0.843637	1.782699	
3	1.856182	-0.624215	-0.766028	-0.733781	0.932968	
4	1.856182	-0.462404	-0.759847	-0.629157	-0.012881	
...	
20635	-0.289187	-0.444985	-0.512592	-0.443449	-1.216128	
20636	-0.845393	-0.888704	-0.944405	-1.008420	-0.691593	
20637	-0.924851	-0.174995	-0.369537	-0.174042	-1.142593	
20638	-0.845393	-0.355600	-0.604429	-0.393753	-1.054583	
20639	-1.004309	0.068408	-0.033977	0.079672	-0.780129	

20640 rows × 6 columns

Handling With Missing Values

```
data.isnull().sum()
```

```
longitude      0
latitude       0
housing_median_age  0
total_rooms    0
total_bedrooms 207
population     0
households     0
median_income  0
median_house_value  0
ocean_proximity  0
dtype: int64
```

```
# here I Will show you imputing values in Null columns only for 'agent' column
data['total_bedrooms'].isnull().sum()
```

207

Simple Imputer

```

from sklearn.impute import SimpleImputer
imputer = SimpleImputer(missing_values=np.nan , strategy='mean')
agent_col = imputer.fit_transform(data['total_bedrooms'].values.reshape(-1,1))
pd.DataFrame(agent_col).isnull().sum()

```

```

0    0
dtype: int64

```

```
data['total_bedrooms'].isnull().sum()
```

```
207
```

Discretization

```

from sklearn.preprocessing import KBinsDiscretizer
temp_data.head()

```

	housing_median_age	total_rooms	population	households	median_income	median_hc
0	41.0	880.0	322.0	126.0	8.3252	
1	21.0	7099.0	2401.0	1138.0	8.3014	
2	52.0	1467.0	496.0	177.0	7.2574	
3	52.0	1274.0	558.0	219.0	5.6431	
4	52.0	1627.0	565.0	259.0	3.8462	

Quantile Discretization Transform

```

trans = KBinsDiscretizer(n_bins =10 , encode = 'ordinal' , strategy='quantile')
new_data = trans.fit_transform(temp_data)
pd.DataFrame(new_data,columns = temp_data.columns )

```

	housing_median_age	total_rooms	population	households	median_income	media
0	8.0	0.0	0.0	0.0	9.0	
1	3.0	9.0	8.0	9.0	9.0	
2	9.0	2.0	0.0	0.0	9.0	

Uniform Discretization Transform

4 9.0 3.0 1.0 2.0 5.0

```
trans = KBinsDiscretizer(n_bins =10 , encode = 'ordinal' , strategy='uniform')
new_data = trans.fit_transform(temp_data)
```

```
pd.DataFrame(new_data,columns = temp_data.columns )
```

	housing_median_age	total_rooms	population	households	median_income	media
0	7.0	0.0	0.0	0.0	5.0	
1	3.0	1.0	0.0	1.0	5.0	
2	9.0	0.0	0.0	0.0	4.0	
3	9.0	0.0	0.0	0.0	3.0	
4	9.0	0.0	0.0	0.0	2.0	
...	
20635	4.0	0.0	0.0	0.0	0.0	
20636	3.0	0.0	0.0	0.0	1.0	
20637	3.0	0.0	0.0	0.0	0.0	
20638	3.0	0.0	0.0	0.0	0.0	
20639	2.0	0.0	0.0	0.0	1.0	

20640 rows × 6 columns

KMeans Discretization Transform

```
trans = KBinsDiscretizer(n_bins =10 , encode = 'ordinal' , strategy='kmeans')
new_data = trans.fit_transform(temp_data)
```

```
pd.DataFrame(new_data,columns = temp_data.columns )
```

	housing_median_age	total_rooms	population	households	median_income	media
0	7.0	0.0	0.0	0.0	6.0	
1	3.0	4.0	3.0	4.0	6.0	
2	9.0	1.0	0.0	0.0	6.0	
3	9.0	0.0	0.0	0.0	4.0	
4	9.0	1.0	0.0	1.0	3.0	
...	
20635	4.0	1.0	1.0	1.0	0.0	
20636	3.0	0.0	0.0	0.0	1.0	
20637	3.0	1.0	1.0	2.0	0.0	
20638	3.0	1.0	0.0	1.0	0.0	
20639	2.0	2.0	1.0	2.0	1.0	

20640 rows × 6 columns